

```
(htrflow_env) (base) apple@LilydeMacBook-Pro-2 htr-comparison-  
project % ./scripts/run_beam_experiment.sh 🐞 开始 Beam Search 参数对比实验 预计总时间: 50-60 分钟  
===== ▶ 测试 num_beams = 1  
===== LilydeMacBook-Pro-2.local - 2025-10-22 22:02:52 UTC -  
INFO - Include schema from 'file:///Users/apple/htrflow_env/lib/python3.11/site-  
packages/xmlschema/schemas/XSD_1.1/xsd11-extra.xsd' LilydeMacBook-Pro-2.local - 2025-10-22  
22:02:52 UTC - ERROR - Could not import OpenMMLab models: []. Traceback (most recent call last):  
File "/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/importer.py", line 26,  
in all_models from htrflow.models.openmmlab.rtmddet import RTMDDet File  
"/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/openmmlab/rtmddet.py", line  
5, in <module> from mmdet.apis import DetInferencer ModuleNotFoundError: No module named 'mmdet'  
LilydeMacBook-Pro-2.local - 2025-10-22 22:02:52 UTC - ERROR - Could not import Teklia models:  
[]. Traceback (most recent call last): File "/Users/apple/htrflow_env/lib/python3.11/site-  
packages/htrflow/models/importer.py", line 37, in all_models from htrflow.models.tekalia.pylaia  
import PyLaia File "/Users/apple/htrflow_env/lib/python3.11/site-  
packages/htrflow/models/tekalia/pylaia.py", line 14, in <module> from laia.common.arguments import  
CommonArgs, DataArgs, DecodeArgs, TrainerArgs ModuleNotFoundError: No module named 'laia'  
LilydeMacBook-Pro-2.local - 2025-10-22 22:02:52 UTC - INFO - Importing 1 input images with batch  
size 1 LilydeMacBook-Pro-2.local - 2025-10-22 22:02:52 UTC - INFO - Initialized collection  
'None' with 1 pages LilydeMacBook-Pro-2.local - 2025-10-22 22:02:52 UTC - INFO - Running step  
Segmentation (step 1 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:02:54 UTC - INFO -  
Initialized YOLO model 'Riksarkivet/yolov9-lines-within-regions-1' from  
/Users/apple/.cache/huggingface/hub/models--Riksarkivet--yolov9-lines-within-regions-  
1/snapshots/f62260d9ae6e782c43c3fd789b7d7e19d64608d9/model.pt on device cpu LilydeMacBook-Pro-  
2.local - 2025-10-22 22:02:54 UTC - INFO - Model 'YOLO' on device 'cpu' received 1 images in  
batches of 1 images per batch (1 batches) YOLO: Running inference (batch size 1): 0%| | 0/1  
[00:00<?, ?it/s] LilydeMacBook-Pro-2.local - 2025-10-22 22:02:54 UTC - INFO - YOLO: Running  
inference on 1 images (batch 1 of 1) YOLO: Running inference (batch size 1): 100%| |  
1/1 [00:02<00:00, 2.25s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:02:56 UTC - INFO - Running  
step TextRecognition (step 2 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:02:57 UTC - INFO -  
Initialized TrOCR model from Riksarkivet/trocr-base-handwritten-hist-swe-2 on device cpu. Using  
a slow image processor as `use_fast` is unset and a slow processor was saved with this model.  
`use_fast=True` will be the default behavior in v4.52, even if the model was saved with a slow  
processor. This will result in minor differences in outputs. You'll still be able to use a slow  
processor with `use_fast=False`. LilydeMacBook-Pro-2.local - 2025-10-22 22:03:00 UTC - INFO -  
Initialized TrOCR processor from Riksarkivet/trocr-base-handwritten-hist-swe-2. LilydeMacBook-  
Pro-2.local - 2025-10-22 22:03:00 UTC - INFO - Model 'TrOCR' on device 'cpu' received 33 images  
in batches of 1 images per batch (33 batches) TrOCR: Running inference (batch size 1): 0%| |  
0/33 [00:00<?, ?it/s] LilydeMacBook-Pro-2.local - 2025-10-22 22:03:00 UTC - INFO - TrOCR: Running  
inference on 1 images (batch 1 of 33) The following generation flags are not valid and may be  
ignored: ['early_stopping', 'length_penalty']. Set `TRANSFORMERS_VERBOSITY=info` for more  
details. TrOCR: Running inference (batch size 1): 3%| | 1/33 [00:14<07:32,
```

14.15s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:03:14 UTC - INFO - TrOCR: Running inference on 1 images (batch 2 of 33) TrOCR: Running inference (batch size 1): 6% | 2/33 [00:32<08:31, 16.49s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:03:32 UTC - INFO - TrOCR: Running inference on 1 images (batch 3 of 33) TrOCR: Running inference (batch size 1): 9% | 3/33 [00:43<07:02, 14.08s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:03:44 UTC - INFO - TrOCR: Running inference on 1 images (batch 4 of 33) TrOCR: Running inference (batch size 1): 12% | 4/33 [00:56<06:32, 13.52s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:03:56 UTC - INFO - TrOCR: Running inference on 1 images (batch 5 of 33) TrOCR: Running inference (batch size 1): 15% | 5/33 [01:09<06:16, 13.46s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:04:10 UTC - INFO - TrOCR: Running inference on 1 images (batch 6 of 33) TrOCR: Running inference (batch size 1): 18% | 6/33 [01:18<05:19, 11.83s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:04:18 UTC - INFO - TrOCR: Running inference on 1 images (batch 7 of 33) TrOCR: Running inference (batch size 1): 21% | 7/33 [01:27<04:47, 11.06s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:04:28 UTC - INFO - TrOCR: Running inference on 1 images (batch 8 of 33) TrOCR: Running inference (batch size 1): 24% | 8/33 [01:36<04:16, 10.25s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:04:36 UTC - INFO - TrOCR: Running inference on 1 images (batch 9 of 33) TrOCR: Running inference (batch size 1): 27% | 9/33 [01:42<03:37, 9.05s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:04:43 UTC - INFO - TrOCR: Running inference on 1 images (batch 10 of 33) TrOCR: Running inference (batch size 1): 30% | 10/33 [01:52<03:34, 9.34s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:04:53 UTC - INFO - TrOCR: Running inference on 1 images (batch 11 of 33) TrOCR: Running inference (batch size 1): 33% | 11/33 [02:03<03:33, 9.68s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:05:03 UTC - INFO - TrOCR: Running inference on 1 images (batch 12 of 33) TrOCR: Running inference (batch size 1): 36% | 12/33 [02:17<03:52, 11.09s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:05:17 UTC - INFO - TrOCR: Running inference on 1 images (batch 13 of 33) TrOCR: Running inference (batch size 1): 39% | 13/33 [02:28<03:39, 10.97s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:05:28 UTC - INFO - TrOCR: Running inference on 1 images (batch 14 of 33) TrOCR: Running inference (batch size 1): 42% | 14/33 [02:39<03:30, 11.10s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:05:40 UTC - INFO - TrOCR: Running inference on 1 images (batch 15 of 33) TrOCR: Running inference (batch size 1): 45% | 15/33 [02:53<03:34, 11.89s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:05:53 UTC - INFO - TrOCR: Running inference on 1 images (batch 16 of 33) TrOCR: Running inference (batch size 1): 48% | 16/33 [03:05<03:26, 12.12s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:06:06 UTC - INFO - TrOCR: Running inference on 1 images (batch 17 of 33) TrOCR: Running inference (batch size 1): 52% | 17/33 [03:18<03:15, 12.23s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:06:18 UTC - INFO - TrOCR: Running inference on 1 images (batch 18 of 33) TrOCR: Running inference (batch size 1): 55% | 18/33 [03:29<02:58, 11.90s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:06:30 UTC - INFO - TrOCR: Running inference on 1 images (batch 19 of 33) TrOCR: Running inference (batch size 1): 58% | 19/33 [03:39<02:38, 11.33s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:06:40 UTC - INFO - TrOCR: Running inference on 1 images (batch 20 of 33) TrOCR: Running inference (batch size 1): 61% | 20/33 [03:51<02:29, 11.52s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:06:51 UTC - INFO - TrOCR: Running inference on 1 images (batch 21 of 33) TrOCR: Running inference (batch size 1): 64% | 21/33 [04:02<02:17,

11.42s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:03 UTC - INFO - TrOCR: Running inference on 1 images (batch 22 of 33) TrOCR: Running inference (batch size 1): 67%|███████| 22/33 [04:04<01:35, 8.65s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:05 UTC - INFO - TrOCR: Running inference on 1 images (batch 23 of 33) TrOCR: Running inference (batch size 1): 70%|███████| 23/33 [04:14<01:29, 8.94s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:14 UTC - INFO - TrOCR: Running inference on 1 images (batch 24 of 33) TrOCR: Running inference (batch size 1): 73%|███████| 24/33 [04:20<01:14, 8.23s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:21 UTC - INFO - TrOCR: Running inference on 1 images (batch 25 of 33) TrOCR: Running inference (batch size 1): 76%|███████| 25/33 [04:30<01:09, 8.74s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:31 UTC - INFO - TrOCR: Running inference on 1 images (batch 26 of 33) TrOCR: Running inference (batch size 1): 79%|███████| 26/33 [04:40<01:02, 8.90s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:40 UTC - INFO - TrOCR: Running inference on 1 images (batch 27 of 33) TrOCR: Running inference (batch size 1): 82%|███████| 27/33 [04:48<00:51, 8.65s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:48 UTC - INFO - TrOCR: Running inference on 1 images (batch 28 of 33) TrOCR: Running inference (batch size 1): 85%|███████| 28/33 [04:49<00:32, 6.52s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:50 UTC - INFO - TrOCR: Running inference on 1 images (batch 29 of 33) TrOCR: Running inference (batch size 1): 88%|███████| 29/33 [04:58<00:29, 7.29s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:07:59 UTC - INFO - TrOCR: Running inference on 1 images (batch 30 of 33) TrOCR: Running inference (batch size 1): 91%|███████| 30/33 [05:06<00:22, 7.48s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:08:07 UTC - INFO - TrOCR: Running inference on 1 images (batch 31 of 33) TrOCR: Running inference (batch size 1): 94%|███████| 31/33 [05:15<00:15, 7.86s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:08:16 UTC - INFO - TrOCR: Running inference on 1 images (batch 32 of 33) TrOCR: Running inference (batch size 1): 97%|███████| 32/33 [05:20<00:06, 6.97s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:08:21 UTC - INFO - TrOCR: Running inference on 1 images (batch 33 of 33) TrOCR: Running inference (batch size 1): 100%|███████| 33/33 [05:22<00:00, 9.76s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:08:22 UTC - INFO - Running step OrderLines (step 3 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:08:22 UTC - INFO - Running step Export (step 4 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:08:22 UTC - INFO - Wrote document to outputs/htrflow/beam1/samples/test.json LilydeMacBook-Pro-2.local - 2025-10-22 22:08:22 UTC - INFO - Processed 1 pages in 329 seconds (average 329.884 seconds per page) ✅ num_beams=1 完成. 用时: 363 秒 🇨🇳 分析结果: ===== 📄 文件分析: test.jpg ===== 📏 图片尺寸: 2776 x 4136 🖼️ 检测行数: 33 🇨🇳 置信度统计: 平均: 0.8289 最高: 0.9731 最低: 0.6462 ⚠️ 低置信度行 (<0.90): Line 0: 0.6551 - Em wy Christimam medh Guds nåde, Sweriges, rätthes ... Line 1: 0.7416 - Karlens Lappgers i Norrlanden uhr Kryarnas och Ede... Line 2: 0.7092 - bohren Hertigimma till Sleßwich Halson, Enkia, Lin... Line 4: 0.7231 - Dom Högborne Fürse och Herre Geer Gers Gustaf Adol... Line 5: 0.6605 - Vttekohöherdhe Kommigh och Arffueste Storfurste ti... Line 6: 0.6462 - fl 3... Line 7: 0.8120 - och Wäissmanneland Wår olskelige Kiäre sonn hafwer... Line 9: 0.8639 - the danske och huruledes hans Hl:r hafwer efterlåt... Line 10: 0.8695 - Elfz tongs Slårt medh någre Härader under så ... Line 13: 0.7244 - och Herres Herr Johans Sweriges, Socher och Wändes... Line 14: 0.8454 - Östergötlande Wår olshelige kiäre Sons Så och Swer... Line 16: 0.8888 - ligh hielp

öfwer heela Ryldet till samma Elfzborgs... Line 17: 0.8674 – göres skulle och wißer och förrogne
män der till f... Line 20: 0.8433 – ohhrlige Edle Wälborna och Wälbördighe Her Abraham... Line
21: 0.8410 – till Nynäs och Skiälighe och Johan Jöranßon till B... Line 22: 0.7169 – bära samma
hielp uthi wår ohrhrlige kiäre yngste J... Line 23: 0.6827 – Freders war Carl Philip Sweriges
Spåhrs och wändes... Line 24: 0.7215 – manneland, Närwahio och Wermeland, frädendome, Sön...
Line 28: 0.8258 – instrubution at utröia på eders en sådan Befasning... Line 30: 0.8808 – händh.
gifwit Giäfle then 22. Septembris Åhr 1613.... Line 32: 0.6833 – Inköpt i Hammarria antoge af
mindre sine 1515... 处理步骤: Segmentation: YOLO TextRecognition: TrOCR

=====

测试 num_beams = 2

LilydeMacBook-Pro-2.local – 2025-10-22 22:08:31 UTC –

INFO – Include schema from 'file:///Users/apple/htrflow_env/lib/python3.11/site-

packages/xmlschema/schemas/XSD_1.1/xsd11-extra.xsd' LilydeMacBook-Pro-2.local – 2025-10-22

22:08:31 UTC – ERROR – Could not import OpenMMLab models: []. Traceback (most recent call last):

File "/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/importer.py", line 26,

in all_models from htrflow.models.openmmlab.rtmDET import RTMDet File

"/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/openmmlab/rtmDET.py", line

5, in <module> from mmdet.apis import DetInferencer ModuleNotFoundError: No module named 'mmdet'

LilydeMacBook-Pro-2.local – 2025-10-22 22:08:31 UTC – ERROR – Could not import Teklia models:

[], Traceback (most recent call last): File "/Users/apple/htrflow_env/lib/python3.11/site-

packages/htrflow/models/importer.py", line 37, in all_models from htrflow.models.tekليا.pyLaia

import PyLaia File "/Users/apple/htrflow_env/lib/python3.11/site-

packages/htrflow/models/tekليا/pyLaia.py", line 14, in <module> from laia.common.arguments import

CommonArgs, DataArgs, DecodeArgs, TrainerArgs ModuleNotFoundError: No module named 'laia'

LilydeMacBook-Pro-2.local – 2025-10-22 22:08:31 UTC – INFO – Importing 1 input images with batch

size 1 LilydeMacBook-Pro-2.local – 2025-10-22 22:08:31 UTC – INFO – Initialized collection

'None' with 1 pages LilydeMacBook-Pro-2.local – 2025-10-22 22:08:31 UTC – INFO – Running step

Segmentation (step 1 / 4) LilydeMacBook-Pro-2.local – 2025-10-22 22:08:32 UTC – INFO –

Initialized YOLO model 'Riksarkivet/yolov9-lines-within-regions-1' from

/Users/apple/.cache/huggingface/hub/models--Riksarkivet--yolov9-lines-within-regions-

1/snapshots/f62260d9ae6e782c43c3fd789b7d7e19d64608d9/model.pt on device cpu LilydeMacBook-Pro-

2.local – 2025-10-22 22:08:32 UTC – INFO – Model 'YOLO' on device 'cpu' received 1 images in

batches of 1 images per batch (1 batches) YOLO: Running inference (batch size 1): 0%| | 0/1

[00:00<?, ?it/s] LilydeMacBook-Pro-2.local – 2025-10-22 22:08:32 UTC – INFO – YOLO: Running

inference on 1 images (batch 1 of 1) YOLO: Running inference (batch size 1): 100%| |

1/1 [00:02<00:00, 2.35s/it] LilydeMacBook-Pro-2.local – 2025-10-22 22:08:34 UTC – INFO – Running

step TextRecognition (step 2 / 4) LilydeMacBook-Pro-2.local – 2025-10-22 22:08:36 UTC – INFO –

Initialized TrOCR model from Riksarkivet/trocr-base-handwritten-hist-swe-2 on device cpu. Using

a slow image processor as 'use_fast' is unset and a slow processor was saved with this model.

'use_fast=True' will be the default behavior in v4.52, even if the model was saved with a slow

processor. This will result in minor differences in outputs. You'll still be able to use a slow

processor with 'use_fast=False'. LilydeMacBook-Pro-2.local – 2025-10-22 22:08:38 UTC – INFO –

Initialized TrOCR processor from Riksarkivet/trocr-base-handwritten-hist-swe-2. LilydeMacBook-Pro-2.local - 2025-10-22 22:08:39 UTC - INFO - Model 'TrOCR' on device 'cpu' received 33 images in batches of 1 images per batch (33 batches) TrOCR: Running inference (batch size 1): 0%| | 0/33 [00:00<?, ?it/s] LilydeMacBook-Pro-2.local - 2025-10-22 22:08:39 UTC - INFO - TrOCR: Running inference on 1 images (batch 1 of 33) TrOCR: Running inference (batch size 1): 3%| | 1/33 [00:13<07:13, 13.55s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:08:52 UTC - INFO - TrOCR: Running inference on 1 images (batch 2 of 33) TrOCR: Running inference (batch size 1): 6%| | 2/33 [00:30<08:07, 15.73s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:09:09 UTC - INFO - TrOCR: Running inference on 1 images (batch 3 of 33) TrOCR: Running inference (batch size 1): 9%| | 3/33 [00:45<07:33, 15.13s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:09:24 UTC - INFO - TrOCR: Running inference on 1 images (batch 4 of 33) TrOCR: Running inference (batch size 1): 12%| | 4/33 [00:57<06:50, 14.15s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:09:36 UTC - INFO - TrOCR: Running inference on 1 images (batch 5 of 33) TrOCR: Running inference (batch size 1): 15%| | 5/33 [01:11<06:26, 13.80s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:09:50 UTC - INFO - TrOCR: Running inference on 1 images (batch 6 of 33) TrOCR: Running inference (batch size 1): 18%| | 6/33 [01:23<06:01, 13.40s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:10:02 UTC - INFO - TrOCR: Running inference on 1 images (batch 7 of 33) TrOCR: Running inference (batch size 1): 21%| | 7/33 [01:41<06:28, 14.92s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:10:20 UTC - INFO - TrOCR: Running inference on 1 images (batch 8 of 33) TrOCR: Running inference (batch size 1): 24%| | 8/33 [01:59<06:32, 15.69s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:10:38 UTC - INFO - TrOCR: Running inference on 1 images (batch 9 of 33) TrOCR: Running inference (batch size 1): 27%| | 9/33 [02:10<05:46, 14.45s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:10:49 UTC - INFO - TrOCR: Running inference on 1 images (batch 10 of 33) TrOCR: Running inference (batch size 1): 30%| | 10/33 [02:32<06:22, 16.61s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:11:11 UTC - INFO - TrOCR: Running inference on 1 images (batch 11 of 33) TrOCR: Running inference (batch size 1): 33%| | 11/33 [02:52<06:27, 17.61s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:11:31 UTC - INFO - TrOCR: Running inference on 1 images (batch 12 of 33) TrOCR: Running inference (batch size 1): 36%| | 12/33 [03:13<06:35, 18.85s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:11:52 UTC - INFO - TrOCR: Running inference on 1 images (batch 13 of 33) TrOCR: Running inference (batch size 1): 39%| | 13/33 [03:31<06:11, 18.58s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:12:10 UTC - INFO - TrOCR: Running inference on 1 images (batch 14 of 33) TrOCR: Running inference (batch size 1): 42%| | 14/33 [03:49<05:46, 18.24s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:12:28 UTC - INFO - TrOCR: Running inference on 1 images (batch 15 of 33) TrOCR: Running inference (batch size 1): 45%| | 15/33 [04:09<05:38, 18.80s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:12:48 UTC - INFO - TrOCR: Running inference on 1 images (batch 16 of 33) TrOCR: Running inference (batch size 1): 48%| | 16/33 [04:27<05:15, 18.55s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:13:06 UTC - INFO - TrOCR: Running inference on 1 images (batch 17 of 33) TrOCR: Running inference (batch size 1): 52%| | 17/33 [04:46<05:00, 18.77s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:13:25 UTC - INFO - TrOCR: Running inference on 1 images (batch 18 of 33) TrOCR: Running inference (batch size 1): 55%| | 18/33 [05:01<04:24, 17.65s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:13:40 UTC - INFO - TrOCR: Running inference on 1 images (batch 19 of 33) TrOCR: Running

inference (batch size 1): 58%|██████| 19/33 [05:17<03:58, 17.02s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:13:56 UTC - INFO - TrOCR: Running inference on 1 images (batch 20 of 33) TrOCR: Running inference (batch size 1): 61%|██████| 20/33 [05:36<03:51, 17.81s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:14:15 UTC - INFO - TrOCR: Running inference on 1 images (batch 21 of 33) TrOCR: Running inference (batch size 1): 64%|██████| 21/33 [05:54<03:32, 17.74s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:14:33 UTC - INFO - TrOCR: Running inference on 1 images (batch 22 of 33) TrOCR: Running inference (batch size 1): 67%|██████| 22/33 [05:58<02:29, 13.58s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:14:37 UTC - INFO - TrOCR: Running inference on 1 images (batch 23 of 33) TrOCR: Running inference (batch size 1): 70%|██████| 23/33 [06:19<02:37, 15.78s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:14:58 UTC - INFO - TrOCR: Running inference on 1 images (batch 24 of 33) TrOCR: Running inference (batch size 1): 73%|██████| 24/33 [06:33<02:18, 15.40s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:15:12 UTC - INFO - TrOCR: Running inference on 1 images (batch 25 of 33) TrOCR: Running inference (batch size 1): 76%|██████| 25/33 [06:55<02:18, 17.31s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:15:34 UTC - INFO - TrOCR: Running inference on 1 images (batch 26 of 33) TrOCR: Running inference (batch size 1): 79%|██████| 26/33 [07:18<02:13, 19.12s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:15:57 UTC - INFO - TrOCR: Running inference on 1 images (batch 27 of 33) TrOCR: Running inference (batch size 1): 82%|██████| 27/33 [07:42<02:03, 20.55s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:16:21 UTC - INFO - TrOCR: Running inference on 1 images (batch 28 of 33) TrOCR: Running inference (batch size 1): 85%|██████| 28/33 [07:46<01:17, 15.42s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:16:25 UTC - INFO - TrOCR: Running inference on 1 images (batch 29 of 33) TrOCR: Running inference (batch size 1): 88%|██████| 29/33 [08:08<01:09, 17.39s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:16:47 UTC - INFO - TrOCR: Running inference on 1 images (batch 30 of 33) TrOCR: Running inference (batch size 1): 91%|██████| 30/33 [08:25<00:52, 17.43s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:17:04 UTC - INFO - TrOCR: Running inference on 1 images (batch 31 of 33) TrOCR: Running inference (batch size 1): 94%|██████| 31/33 [08:45<00:36, 18.06s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:17:24 UTC - INFO - TrOCR: Running inference on 1 images (batch 32 of 33) TrOCR: Running inference (batch size 1): 97%|██████| 32/33 [08:55<00:15, 15.78s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:17:34 UTC - INFO - TrOCR: Running inference on 1 images (batch 33 of 33) TrOCR: Running inference (batch size 1): 100%|██████| 33/33 [08:58<00:00, 16.32s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:17:37 UTC - INFO - Running step OrderLines (step 3 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:17:37 UTC - INFO - Running step Export (step 4 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:17:37 UTC - INFO - Wrote document to outputs/htrflow/beam2/samples/test.json LilydeMacBook-Pro-2.local - 2025-10-22 22:17:37 UTC - INFO - Processed 1 pages in 546 seconds (average 546.467 seconds per page) ✅
num_beams=2 完成, 用时: 555 秒 🇨🇳 分析结果:

===== 📄 文件分析: test.jpg
===== 📏 图片尺寸: 2776 x 4136 📄 检测行数:
33 🇨🇳 置信度统计: 平均: 0.9882 最高: 0.9991 最低: 0.8645 ⚠️ 低置信度行 (<0.90): Line 6: 0.8645 - fl 3...
🔪 处理步骤: Segmentation: YOLO TextRecognition: TrOCR
=====

测试 num_beams = 4

LilydeMacBook-Pro-2.local - 2025-10-22 22:17:47 UTC -

INFO - Include schema from 'file:///Users/apple/htrflow_env/lib/python3.11/site-packages/xmlschema/schemas/XSD_1.1/xsd11-extra.xsd' LilydeMacBook-Pro-2.local - 2025-10-22 22:17:47 UTC - ERROR - Could not import OpenMMLab models: []. Traceback (most recent call last): File "/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/importer.py", line 26, in all_models from htrflow.models.openmmlab.rtmddet import RTMDDet File "/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/openmmlab/rtmddet.py", line 5, in <module> from mmdet.apis import DetInferencer ModuleNotFoundError: No module named 'mmdet' LilydeMacBook-Pro-2.local - 2025-10-22 22:17:47 UTC - ERROR - Could not import Teklia models: []. Traceback (most recent call last): File "/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/importer.py", line 37, in all_models from htrflow.models.tekalia.pylaia import PyLaia File "/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/tekalia/pylaia.py", line 14, in <module> from laia.common.arguments import CommonArgs, DataArgs, DecodeArgs, TrainerArgs ModuleNotFoundError: No module named 'laia' LilydeMacBook-Pro-2.local - 2025-10-22 22:17:47 UTC - INFO - Importing 1 input images with batch size 1 LilydeMacBook-Pro-2.local - 2025-10-22 22:17:48 UTC - INFO - Initialized collection 'None' with 1 pages LilydeMacBook-Pro-2.local - 2025-10-22 22:17:48 UTC - INFO - Running step Segmentation (step 1 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:17:49 UTC - INFO - Initialized YOLO model 'Riksarkivet/yolov9-lines-within-regions-1' from /Users/apple/.cache/huggingface/hub/models--Riksarkivet--yolov9-lines-within-regions-1/snapshots/f62260d9ae6e782c43c3fd789b7d7e19d64608d9/model.pt on device cpu LilydeMacBook-Pro-2.local - 2025-10-22 22:17:49 UTC - INFO - Model 'YOLO' on device 'cpu' received 1 images in batches of 1 images per batch (1 batches) YOLO: Running inference (batch size 1): 0%| | 0/1 [00:00<?, ?it/s] LilydeMacBook-Pro-2.local - 2025-10-22 22:17:49 UTC - INFO - YOLO: Running inference on 1 images (batch 1 of 1) YOLO: Running inference (batch size 1): 100%| | 1/1 [00:02<00:00, 2.45s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:17:51 UTC - INFO - Running step TextRecognition (step 2 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:17:53 UTC - INFO - Initialized TrOCR model from Riksarkivet/trocr-base-handwritten-hist-swe-2 on device cpu. Using a slow image processor as 'use_fast' is unset and a slow processor was saved with this model. 'use_fast=True' will be the default behavior in v4.52, even if the model was saved with a slow processor. This will result in minor differences in outputs. You'll still be able to use a slow processor with 'use_fast=False'. LilydeMacBook-Pro-2.local - 2025-10-22 22:17:57 UTC - INFO - Initialized TrOCR processor from Riksarkivet/trocr-base-handwritten-hist-swe-2. LilydeMacBook-Pro-2.local - 2025-10-22 22:17:58 UTC - INFO - Model 'TrOCR' on device 'cpu' received 33 images in batches of 1 images per batch (33 batches) TrOCR: Running inference (batch size 1): 0%| | 0/33 [00:00<?, ?it/s] LilydeMacBook-Pro-2.local - 2025-10-22 22:17:58 UTC - INFO - TrOCR: Running inference on 1 images (batch 1 of 33) TrOCR: Running inference (batch size 1): 3%| | 1/33 [00:44<23:37, 44.31s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:18:42 UTC - INFO - TrOCR: Running inference on 1 images (batch 2 of 33) TrOCR: Running inference (batch size 1): 6%| | 2/33 [01:24<21:42, 42.01s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:19:22 UTC - INFO - TrOCR: Running inference on 1 images (batch 3 of 33) TrOCR: Running inference (batch size 1): 9%| |

3/33 [02:02<19:56, 39.90s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:20:00 UTC - INFO - TrOCR: Running inference on 1 images (batch 4 of 33) TrOCR: Running inference (batch size 1): 12% |

4/33 [02:35<18:05, 37.42s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:20:33 UTC - INFO - TrOCR: Running inference on 1 images (batch 5 of 33) TrOCR: Running inference (batch size 1): 15% |

5/33 [03:15<17:47, 38.13s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:21:13 UTC - INFO - TrOCR: Running inference on 1 images (batch 6 of 33) TrOCR: Running inference (batch size 1): 18% |

6/33 [03:48<16:28, 36.61s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:21:46 UTC - INFO - TrOCR: Running inference on 1 images (batch 7 of 33) TrOCR: Running inference (batch size 1): 21% |

7/33 [04:21<15:18, 35.32s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:22:19 UTC - INFO - TrOCR: Running inference on 1 images (batch 8 of 33) TrOCR: Running inference (batch size 1): 24% |

8/33 [04:51<14:03, 33.75s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:22:49 UTC - INFO - TrOCR: Running inference on 1 images (batch 9 of 33) TrOCR: Running inference (batch size 1): 27% |

9/33 [05:09<11:30, 28.76s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:23:07 UTC - INFO - TrOCR: Running inference on 1 images (batch 10 of 33) TrOCR: Running inference (batch size 1): 30% |

10/33 [05:38<11:05, 28.92s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:23:36 UTC - INFO - TrOCR: Running inference on 1 images (batch 11 of 33) TrOCR: Running inference (batch size 1): 33% |

11/33 [06:04<10:16, 28.02s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:24:02 UTC - INFO - TrOCR: Running inference on 1 images (batch 12 of 33) TrOCR: Running inference (batch size 1): 36% |

12/33 [06:34<10:01, 28.63s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:24:32 UTC - INFO - TrOCR: Running inference on 1 images (batch 13 of 33) TrOCR: Running inference (batch size 1): 39% |

13/33 [06:59<09:09, 27.46s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:24:57 UTC - INFO - TrOCR: Running inference on 1 images (batch 14 of 33) TrOCR: Running inference (batch size 1): 42% |

14/33 [07:22<08:16, 26.12s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:25:20 UTC - INFO - TrOCR: Running inference on 1 images (batch 15 of 33) TrOCR: Running inference (batch size 1): 45% |

15/33 [07:49<07:56, 26.49s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:25:48 UTC - INFO - TrOCR: Running inference on 1 images (batch 16 of 33) TrOCR: Running inference (batch size 1): 48% |

16/33 [08:13<07:14, 25.56s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:26:11 UTC - INFO - TrOCR: Running inference on 1 images (batch 17 of 33) TrOCR: Running inference (batch size 1): 52% |

17/33 [08:37<06:42, 25.13s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:26:35 UTC - INFO - TrOCR: Running inference on 1 images (batch 18 of 33) TrOCR: Running inference (batch size 1): 55% |

18/33 [08:58<05:56, 23.77s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:26:56 UTC - INFO - TrOCR: Running inference on 1 images (batch 19 of 33) TrOCR: Running inference (batch size 1): 58% |

19/33 [09:19<05:21, 22.99s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:27:17 UTC - INFO - TrOCR: Running inference on 1 images (batch 20 of 33) TrOCR: Running inference (batch size 1): 61% |

20/33 [09:46<05:13, 24.14s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:27:44 UTC - INFO - TrOCR: Running inference on 1 images (batch 21 of 33) TrOCR: Running inference (batch size 1): 64% |

21/33 [10:10<04:50, 24.23s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:28:08 UTC - INFO - TrOCR: Running inference on 1 images (batch 22 of 33) TrOCR: Running inference (batch size 1): 67% |

22/33 [10:15<03:22, 18.41s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:28:13 UTC - INFO - TrOCR: Running inference on 1 images (batch 23 of 33) TrOCR: Running inference (batch size 1):

70%|██████| 23/33 [10:43<03:34, 21.42s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:28:41 UTC
- INFO - TrOCR: Running inference on 1 images (batch 24 of 33) TrOCR: Running inference (batch
size 1): 73%|██████| 24/33 [11:03<03:07, 20.85s/it] LilydeMacBook-Pro-2.local - 2025-10-22
22:29:01 UTC - INFO - TrOCR: Running inference on 1 images (batch 25 of 33) TrOCR: Running
inference (batch size 1): 76%|██████| 25/33 [11:34<03:11, 23.91s/it] LilydeMacBook-Pro-2.local
- 2025-10-22 22:29:32 UTC - INFO - TrOCR: Running inference on 1 images (batch 26 of 33) TrOCR:
Running inference (batch size 1): 79%|██████| 26/33 [12:00<02:50, 24.42s/it] LilydeMacBook-
Pro-2.local - 2025-10-22 22:29:58 UTC - INFO - TrOCR: Running inference on 1 images (batch 27 of
33) TrOCR: Running inference (batch size 1): 82%|██████| 27/33 [12:25<02:28,
24.80s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:30:23 UTC - INFO - TrOCR: Running inference
on 1 images (batch 28 of 33) TrOCR: Running inference (batch size 1): 85%|██████| 28/33
[12:28<01:31, 18.28s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:30:26 UTC - INFO - TrOCR:
Running inference on 1 images (batch 29 of 33) TrOCR: Running inference (batch size 1):
88%|██████| 29/33 [12:55<01:23, 20.93s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:30:53
UTC - INFO - TrOCR: Running inference on 1 images (batch 30 of 33) TrOCR: Running inference
(batch size 1): 91%|██████| 30/33 [13:19<01:05, 21.76s/it] LilydeMacBook-Pro-2.local - 2025-
10-22 22:31:17 UTC - INFO - TrOCR: Running inference on 1 images (batch 31 of 33) TrOCR: Running
inference (batch size 1): 94%|██████| 31/33 [13:48<00:47, 23.97s/it] LilydeMacBook-Pro-
2.local - 2025-10-22 22:31:46 UTC - INFO - TrOCR: Running inference on 1 images (batch 32 of 33)
TrOCR: Running inference (batch size 1): 97%|██████| 32/33 [14:04<00:21,
21.58s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:32:02 UTC - INFO - TrOCR: Running inference
on 1 images (batch 33 of 33) TrOCR: Running inference (batch size 1): 100%|██████| 33/33
[14:07<00:00, 25.68s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:32:05 UTC - INFO - Running
step OrderLines (step 3 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:32:05 UTC - INFO -
Running step Export (step 4 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:32:05 UTC - INFO -
Wrote document to outputs/htrflow/beam4/samples/test.json LilydeMacBook-Pro-2.local - 2025-10-22
22:32:05 UTC - INFO - Processed 1 pages in 857 seconds (average 857.770 seconds per page) ✓
num_beams=4 完成, 用时: 868 秒 🇨🇳 分析结果:

===== 📄 文件分析: test.jpg
===== 📏 图片尺寸: 2776 x 4136 📄 检测行数:
33 🇨🇳 置信度统计: 平均: 0.9617 最高: 0.9979 最低: 0.1397 ⚠️ 低置信度行 (<0.90): Line 3: 0.1397 - No
7... Line 6: 0.8645 - fl 3... 🛠️ 处理步骤: Segmentation: YOLO TextRecognition: TrOCR

=====
===== ▶️ 测试 num_beams = 8
===== LilydeMacBook-Pro-2.local - 2025-10-22 22:32:14 UTC -
INFO - Include schema from 'file:///Users/apple/htrflow_env/lib/python3.11/site-
packages/xmlschema/schemas/XSD_1.1/xsd11-extra.xsd' LilydeMacBook-Pro-2.local - 2025-10-22
22:32:14 UTC - ERROR - Could not import OpenMMLab models: []. Traceback (most recent call last):
File "/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/importer.py", line 26,
in all_models from htrflow.models.openmmlab.rtmddet import RTMDDet File
"/Users/apple/htrflow_env/lib/python3.11/site-packages/htrflow/models/openmmlab/rtmddet.py", line
5, in <module> from mmdet.apis import DetInferencer ModuleNotFoundError: No module named 'mmdet'

```
LilydeMacBook-Pro-2.local - 2025-10-22 22:32:14 UTC - ERROR - Could not import Teklia models:
[. Traceback (most recent call last): File "/Users/apple/htrflow_env/lib/python3.11/site-
packages/htrflow/models/importer.py", line 37, in all_models from htrflow.models.teklia.pylaia
import PyLaia File "/Users/apple/htrflow_env/lib/python3.11/site-
packages/htrflow/models/teklia/pylaia.py", line 14, in <module> from laia.common.arguments import
CommonArgs, DataArgs, DecodeArgs, TrainerArgs ModuleNotFoundError: No module named 'laia'
LilydeMacBook-Pro-2.local - 2025-10-22 22:32:14 UTC - INFO - Importing 1 input images with batch
size 1 LilydeMacBook-Pro-2.local - 2025-10-22 22:32:14 UTC - INFO - Initialized collection
'None' with 1 pages LilydeMacBook-Pro-2.local - 2025-10-22 22:32:14 UTC - INFO - Running step
Segmentation (step 1 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:32:15 UTC - INFO -
Initialized YOLO model 'Riksarkivet/yolov9-lines-within-regions-1' from
/Users/apple/.cache/huggingface/hub/models--Riksarkivet--yolov9-lines-within-regions-
1/snapshots/f62260d9ae6e782c43c3fd789b7d7e19d64608d9/model.pt on device cpu LilydeMacBook-Pro-
2.local - 2025-10-22 22:32:15 UTC - INFO - Model 'YOLO' on device 'cpu' received 1 images in
batches of 1 images per batch (1 batches) YOLO: Running inference (batch size 1): 0%| | 0/1
[00:00<?, ?it/s] LilydeMacBook-Pro-2.local - 2025-10-22 22:32:15 UTC - INFO - YOLO: Running
inference on 1 images (batch 1 of 1) YOLO: Running inference (batch size 1): 100%| | 1/1
[00:01<00:00, 2.00s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:32:17 UTC - INFO - Running
step TextRecognition (step 2 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:32:19 UTC - INFO -
Initialized TrOCR model from Riksarkivet/trocr-base-handwritten-hist-swe-2 on device cpu. Using
a slow image processor as `use_fast` is unset and a slow processor was saved with this model.
`use_fast=True` will be the default behavior in v4.52, even if the model was saved with a slow
processor. This will result in minor differences in outputs. You'll still be able to use a slow
processor with `use_fast=False`. LilydeMacBook-Pro-2.local - 2025-10-22 22:32:21 UTC - INFO -
Initialized TrOCR processor from Riksarkivet/trocr-base-handwritten-hist-swe-2. LilydeMacBook-
Pro-2.local - 2025-10-22 22:32:22 UTC - INFO - Model 'TrOCR' on device 'cpu' received 33 images
in batches of 1 images per batch (33 batches) TrOCR: Running inference (batch size 1): 0%| |
0/33 [00:00<?, ?it/s] LilydeMacBook-Pro-2.local - 2025-10-22 22:32:23 UTC - INFO - TrOCR: Running
inference on 1 images (batch 1 of 33) TrOCR: Running inference (batch size 1): 3%| | 1/33
[00:50<26:50, 50.33s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:33:13 UTC - INFO - TrOCR:
Running inference on 1 images (batch 2 of 33) TrOCR: Running inference (batch size 1): 6%| | 2/33
[01:42<26:26, 51.18s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:34:05 UTC - INFO - TrOCR:
Running inference on 1 images (batch 3 of 33) TrOCR: Running inference (batch size 1): 9%| | 3/33
[02:28<24:31, 49.05s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:34:51 UTC - INFO - TrOCR:
Running inference on 1 images (batch 4 of 33) TrOCR: Running inference (batch size 1): 12%| | 4/33
[03:10<22:14, 46.03s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:35:32 UTC - INFO - TrOCR:
Running inference on 1 images (batch 5 of 33) TrOCR: Running inference (batch size 1): 15%| | 5/33
[04:03<22:39, 48.54s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:36:25 UTC - INFO - TrOCR:
Running inference on 1 images (batch 6 of 33) TrOCR: Running inference (batch size 1): 18%| | 6/33
[04:51<21:48, 48.47s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:37:14 UTC - INFO - TrOCR:
Running inference on 1 images (batch 7 of 33) TrOCR: Running inference (batch size 1): 21%| | 7/33
[05:38<20:49, 48.07s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:38:01 UTC - INFO - TrOCR:
```

Running inference on 1 images (batch 8 of 33) TrOCR: Running inference (batch size 1): 24% |
8/33 [06:26<20:03, 48.15s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:38:49 UTC - INFO - TrOCR:
Running inference on 1 images (batch 9 of 33) TrOCR: Running inference (batch size 1): 27% |
9/33 [06:57<17:02, 42.59s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:39:20 UTC - INFO - TrOCR:
Running inference on 1 images (batch 10 of 33) TrOCR: Running inference (batch size 1): 30% |
10/33 [07:50<17:34, 45.86s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:40:13 UTC - INFO -
TrOCR: Running inference on 1 images (batch 11 of 33) TrOCR: Running inference (batch size 1):
33% | 11/33 [08:37<16:59, 46.35s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:41:00 UTC -
INFO - TrOCR: Running inference on 1 images (batch 12 of 33) TrOCR: Running inference (batch
size 1): 36% | 12/33 [09:36<17:31, 50.06s/it]LilydeMacBook-Pro-2.local - 2025-10-22
22:41:59 UTC - INFO - TrOCR: Running inference on 1 images (batch 13 of 33) TrOCR: Running
inference (batch size 1): 39% | 13/33 [10:21<16:12, 48.63s/it]LilydeMacBook-Pro-2.local -
2025-10-22 22:42:44 UTC - INFO - TrOCR: Running inference on 1 images (batch 14 of 33) TrOCR:
Running inference (batch size 1): 42% | 14/33 [11:05<14:53, 47.01s/it]LilydeMacBook-Pro-
2.local - 2025-10-22 22:43:28 UTC - INFO - TrOCR: Running inference on 1 images (batch 15 of 33)
TrOCR: Running inference (batch size 1): 45% | 15/33 [11:56<14:27,
48.21s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:44:19 UTC - INFO - TrOCR: Running inference
on 1 images (batch 16 of 33) TrOCR: Running inference (batch size 1): 48% | 16/33
[12:42<13:28, 47.59s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:45:05 UTC - INFO - TrOCR:
Running inference on 1 images (batch 17 of 33) TrOCR: Running inference (batch size 1):
52% | 17/33 [13:26<12:25, 46.61s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:45:49 UTC -
INFO - TrOCR: Running inference on 1 images (batch 18 of 33) TrOCR: Running inference (batch
size 1): 55% | 18/33 [14:05<11:03, 44.23s/it]LilydeMacBook-Pro-2.local - 2025-10-22
22:46:28 UTC - INFO - TrOCR: Running inference on 1 images (batch 19 of 33) TrOCR: Running
inference (batch size 1): 58% | 19/33 [14:46<10:04, 43.21s/it]LilydeMacBook-Pro-2.local -
2025-10-22 22:47:09 UTC - INFO - TrOCR: Running inference on 1 images (batch 20 of 33) TrOCR:
Running inference (batch size 1): 61% | 20/33 [15:41<10:10, 46.97s/it]LilydeMacBook-Pro-
2.local - 2025-10-22 22:48:04 UTC - INFO - TrOCR: Running inference on 1 images (batch 21 of 33)
TrOCR: Running inference (batch size 1): 64% | 21/33 [16:32<09:38,
48.23s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:48:55 UTC - INFO - TrOCR: Running inference
on 1 images (batch 22 of 33) TrOCR: Running inference (batch size 1): 67% | 22/33
[16:43<06:44, 36.78s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:49:06 UTC - INFO - TrOCR:
Running inference on 1 images (batch 23 of 33) TrOCR: Running inference (batch size 1):
70% | 23/33 [17:50<07:39, 45.99s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:50:13 UTC
- INFO - TrOCR: Running inference on 1 images (batch 24 of 33) TrOCR: Running inference (batch
size 1): 73% | 24/33 [18:38<06:58, 46.47s/it]LilydeMacBook-Pro-2.local - 2025-10-22
22:51:01 UTC - INFO - TrOCR: Running inference on 1 images (batch 25 of 33) TrOCR: Running
inference (batch size 1): 76% | 25/33 [19:54<07:23, 55.49s/it]LilydeMacBook-Pro-2.local
- 2025-10-22 22:52:17 UTC - INFO - TrOCR: Running inference on 1 images (batch 26 of 33) TrOCR:
Running inference (batch size 1): 79% | 26/33 [20:40<06:08, 52.66s/it]LilydeMacBook-
Pro-2.local - 2025-10-22 22:53:03 UTC - INFO - TrOCR: Running inference on 1 images (batch 27 of
33) TrOCR: Running inference (batch size 1): 82% | 27/33 [21:26<05:03,

```
50.53s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:53:49 UTC - INFO - TrOCR: Running inference
on 1 images (batch 28 of 33) TrOCR: Running inference (batch size 1): 85%|██████████| 28/33
[21:31<03:05, 37.04s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:53:54 UTC - INFO - TrOCR:
Running inference on 1 images (batch 29 of 33) TrOCR: Running inference (batch size 1):
88%|██████████| 29/33 [22:21<02:43, 40.76s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:54:44
UTC - INFO - TrOCR: Running inference on 1 images (batch 30 of 33) TrOCR: Running inference
(batch size 1): 91%|██████████| 30/33 [23:11<02:10, 43.62s/it]LilydeMacBook-Pro-2.local - 2025-
10-22 22:55:34 UTC - INFO - TrOCR: Running inference on 1 images (batch 31 of 33) TrOCR: Running
inference (batch size 1): 94%|██████████| 31/33 [24:05<01:33, 46.67s/it]LilydeMacBook-Pro-
2.local - 2025-10-22 22:56:28 UTC - INFO - TrOCR: Running inference on 1 images (batch 32 of 33)
TrOCR: Running inference (batch size 1): 97%|██████████| 32/33 [24:29<00:39,
39.92s/it]LilydeMacBook-Pro-2.local - 2025-10-22 22:56:52 UTC - INFO - TrOCR: Running inference
on 1 images (batch 33 of 33) TrOCR: Running inference (batch size 1): 100%|██████████| 33/33
[24:34<00:00, 44.69s/it] LilydeMacBook-Pro-2.local - 2025-10-22 22:56:57 UTC - INFO - Running
step OrderLines (step 3 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:56:57 UTC - INFO -
Running step Export (step 4 / 4) LilydeMacBook-Pro-2.local - 2025-10-22 22:56:57 UTC - INFO -
Wrote document to outputs/htrflow/beam8/samples/test.json LilydeMacBook-Pro-2.local - 2025-10-22
22:56:57 UTC - INFO - Processed 1 pages in 1483 seconds (average 1483.179 seconds per page) ✓
num_beams=8 完成, 用时: 1492 秒 🇨🇳 分析结果:
===== 📄 文件分析: test.jpg
===== 📏 图片尺寸: 2776 x 4136 📄 检测行数:
33 🇨🇳 置信度统计: 平均: 0.9882 最高: 0.9985 最低: 0.8645 ⚠️ 低置信度行 (<0.90): Line 6: 0.8645 - fl 3...
🔧 处理步骤: Segmentation: YOLO TextRecognition: TrOCR
===== 🎉 所有实验完成! 查看结果: ls -lh
outputs/htrflow/beam*/
```